

**Московский государственный технический
университет им. Н. Э. Баумана**
Факультет «Информатика и системы управления»

Кафедра «Системы обработки информации и управления»
Курс «Технологии машинного обучения»

Отчет по лабораторной работе №5
Линейные модели, SVM и деревья решений

Группа: РТ5-61
Студент: Савушкин Д. А.
Преподаватель: Гапанюк Ю. Е.

Москва, 2020 г.

Цель лабораторной работы: изучение линейных моделей, SVM и деревьев решений.

Задание:

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите следующие модели:
 - 4.1. одну из линейных моделей;
 - 4.2. SVM;
 - 4.3. дерево решений.
5. Оцените качество моделей с помощью двух подходящих для задачи метрик. Сравните качество полученных моделей.

Дополнительные задания:

1. Проведите эксперименты с важностью признаков в дереве решений.
2. Визуализируйте дерево решений.

Текст программы и экранные формы с примерами выполнения программы:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.impute import SimpleImputer
import pandas_profiling as pp
import warnings
warnings.simplefilter("ignore")
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
import pandas.util.testing as tm

[ ] data = pd.read_csv('datasets_494724_1288143_COVID19_line_list_data.csv')
data.head()

[ ] is_in_country reporting Unnamed: 3 summary location country gender age symptom_onset If_onset_approximated hosp_visit_date exposure_start exposure_end visiting Wuhan from Wuhan death recovered symptom source
First confirmed imported COVID-19 pneumonia pa...
NaN 1/20/2020 NaN Shenzhen, Guangdong China male 66.0 01/03/20 0.0 01/11/20 12/29/2019 01/04/20 1 0.0 0 0 NaN Shenzhen Municipal Health Commission http://wjw.sz.gov.cn
First confirmed imported COVID-19 pneumonia pa...
NaN 1/20/2020 NaN Shanghai China female 56.0 1/15/2020 0.0 1/15/2020 NaN 01/12/20 0 1.0 0 0 NaN Official Weibo of Shanghai Municipal Health Co... https://www.we...
First confirmed imported cases in Zhejiang. pa...
NaN 1/21/2020 NaN Zhejiang China male 46.0 01/04/20 0.0 1/17/2020 NaN 01/03/20 0 1.0 0 0 NaN Health Commission of Zhejiang Province http://www.zjzf.gov.cn
new confirmed imported COVID-19 pneumonia in T...
NaN 1/21/2020 NaN Tianjin China female 60.0 NaN NaN 1/19/2020 NaN NaN 1 0.0 0 0 NaN 人民日报官方微博 https://m.weibo.com
```

```
[ ] # проверим есть ли пропущенные значения
data.isnull().sum()

[ ] id           0
case_in_country 197
reporting_date   1
Unnamed: 3       1085
summary          5
location          0
country           0
gender            183
age               242
symptom_onset    522
if_onset_approximated 525
hosp_visit_date 578
exposure_start   957
exposure_end     744
visiting_Wuhan   0
from_Wuhan        4
death              0
recovered          0
symptom           815
source             0
link              0
Unnamed: 21      1085
Unnamed: 22      1085
Unnamed: 23      1085
Unnamed: 24      1085
Unnamed: 25      1085
Unnamed: 26      1085
dtype: int64
```

```
[ ] # Удаление колонок, содержащих пустые значения
data_new_1 = data.drop(columns=['Unnamed: 21','link','source', 'summary','reporting_date','Unnamed: 22','Unnamed: 23','Unnamed: 24','Unnamed: 25','symptom','Unnamed: 26', 'Unnamed: 3','symptom_onset','if_onset_approximated','hosp_visit_date'])
(data.shape, data_new_1.shape)

[ ] ((1085, 27), (1085, 10))
```

```
[ ] data_new_1.head()
```

	id	case_in_country	location	country	gender	age	visiting_Wuhan	from_Wuhan	death	recovered
0	1	NaN	Shenzhen, Guangdong	China	male	66.0	1	0.0	0	0
1	2	NaN	Shanghai	China	female	56.0	0	1.0	0	0
2	3	NaN	Zhailian, China	China	male	46.0	0	1.0	0	0

```
[ ] data_new_1.isnull().sum()
```

```
[ ] id           0
case_in_country 197
location          0
country           0
gender            183
age               242
visiting_Wuhan   0
from_Wuhan        4
death              0
recovered          0
dtype: int64
```

```
[ ] # удаление строки, содержащих пустые значения
data_new_2 = data_new_1.dropna(axis=0,how='any',subset=['case_in_country','gender', 'age'])

(data_new_1.shape, data_new_2.shape)
```

```
[ ] ((1085, 10), (635, 10))
```

```
[ ] data_new_2.head()
```

	id	case_in_country	location	country	gender	age	visiting_Wuhan	from_Wuhan	death	recovered
197	198	1.0	Bordeaux	France	male	48.0	1	0.0	0	0
198	199	2.0	Paris	France	male	31.0	0	1.0	0	02/12/20
199	200	3.0	Paris	France	female	30.0	0	1.0	0	02/12/20
200	201	4.0	Paris	France	male	80.0	0	1.0	2/14/2020	0
209	210	13.0	Paris	France	female	33.0	0	0.0	0	0

```
[ ] data_new_2.shape[0]
```

```
[ ] 635
```

```
[ ] data_new_2['death'] = data_new_2['death'].apply(lambda x: 0 if x=='0' else 1)
```

```
[ ] data_new_2['recovered'] = data_new_2['recovered'].apply(lambda x: 0 if x=='0' else 1)
```

data_new_2

197	198	1.0	Bordeaux	France	male	48.0		1	0.0	0	0	0
198	199	2.0	Paris	France	male	31.0		0	1.0	0	0	1
199	200	3.0	Paris	France	female	30.0		0	1.0	0	0	1
200	201	4.0	Paris	France	male	80.0		0	1.0	1	0	0
209	210	13.0	Paris	France	female	33.0		0	0.0	0	0	0
...
1027	1028	32.0	Andalusia	Spain	male	58.0		0	0.0	0	0	0
1029	1030	34.0	Zaragoza	Spain	female	27.0		0	0.0	0	0	0
1030	1031	1.0	Jonkoping	Sweden	female	25.0		1	0.0	0	0	0
1052	1053	1.0	Lebanon	Lebanon	female	45.0		0	0.0	0	0	0
1084	1085	1.0	Bern	Switzerland	male	70.0		0	0.0	0	0	0

635 rows × 10 columns

```
[ ] data = pd.get_dummies(data_new_2)
data = data.drop(columns=['id','case_in_country'])
(data.shape, data.shape)
```

(635, 118), (635, 118)

```
[ ] from sklearn.model_selection import train_test_split
```

```
[ ] y = data.death
data.drop('death', axis=1, inplace=True)
X_cov, X_test, y_cov, y_test = train_test_split(data, y, test_size=0.2)
print (X_cov.shape, y_cov.shape)
print (X_test.shape, y_test.shape)
```

(508, 117) (508,)
(127, 117) (127,)

Обучение модели

```
[ ] from sklearn.neighbors import KNeighborsClassifier  
  
[ ] KNeighborsClassifierObj = KNeighborsClassifier(n_neighbors=10)  
  
[ ] KNeighborsClassifierObj.fit(X_cov, y_cov)  
  
👤 KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
    metric_params=None, n_jobs=None, n_neighbors=10, p=2,  
    weights='uniform')  
  
[ ] y_predicted = KNeighborsClassifierObj.predict(X_test)
```

Метрика качества

```
[ ] from sklearn.metrics import accuracy_score, balanced_accuracy_score, precision_score, f1_score, classification_report  
  
[ ] accuracy_score(y_test, y_predicted)  
  
👤 0.9606299212598425  
  
[ ] precision_score(y_test, y_predicted)  
  
👤 0.0
```

Смертность не зависит от других параметров.

```
[ ] classification_report(y_test, y_predicted, output_dict = True)

❸ {'0': {'f1-score': 0.9799196787148594,
         'precision': 0.9606299212598425,
         'recall': 1.0,
         'support': 122},
     '1': {'f1-score': 0.0, 'precision': 0.0, 'recall': 0.0, 'support': 5},
     'accuracy': 0.9606299212598425,
     'macro avg': {'f1-score': 0.4899598393574297,
                    'precision': 0.48031496062992124,
                    'recall': 0.5,
                    'support': 127},
     'weighted avg': {'f1-score': 0.9413401638048255,
                      'precision': 0.9228098456196913,
                      'recall': 0.9606299212598425,
                      'support': 127}}
```

Кросс-валидация

```
[ ] from sklearn.model_selection import cross_val_score

[ ] scores = cross_val_score(KNeighborsClassifierObj,
                             X_cov, y_cov, cv=3,
                             scoring='f1_weighted')
scores, np.mean(scores)

❸ (array([0.95610184, 0.96463877, 0.95584342]), 0.9588613436340419)
```

Подбор гиперпараметров

```
[ ] from sklearn.model_selection import GridSearchCV  
  
[ ] n_range = np.array(range(5,55,5))  
tuned_parameters = [{'n_neighbors': n_range}  
  
[ ] clf_gs = GridSearchCV(KNeighborsClassifier(), tuned_parameters, cv=5, scoring='f1_weighted')  
  
[ ] clf_gs.fit(X_cov, y_cov)  
  
GridSearchCV(cv=5, error_score=nan,  
             estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,  
                                             metric='minkowski',  
                                             metric_params=None, n_jobs=None,  
                                             n_neighbors=5, p=2,  
                                             weights='uniform'),  
             iid='deprecated', n_jobs=None,  
             param_grid=[{'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50])}],  
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
             scoring='f1_weighted', verbose=0)
```

```
[ ] clf_gs.best_params_
```

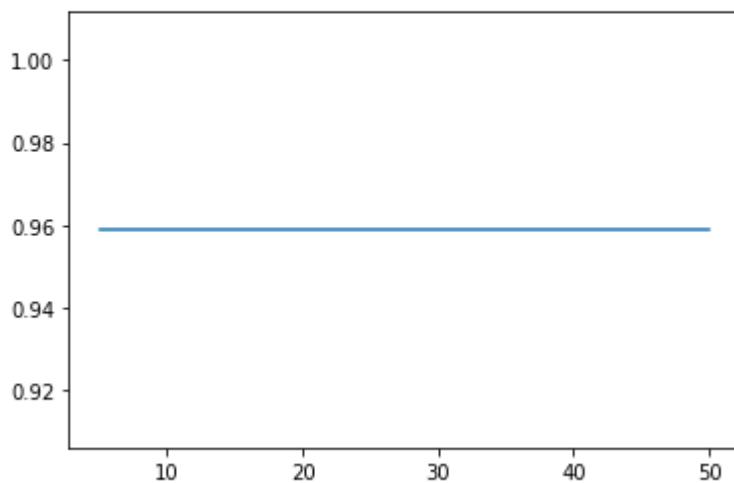
```
{'n_neighbors': 5}
```

```
[ ] clf_gs.best_score_
```

```
0.9588742014124028
```

```
[ ] plt.plot(n_range, clf_gs.cv_results_['mean_test_score'])
```

```
[<matplotlib.lines.Line2D at 0x7f1752b97278>]
```



Логическая регрессия

```
[ ] from sklearn.linear_model import LogisticRegression

[ ] LogRegression = LogisticRegression()

[ ] LogRegression.fit(X_cov,y_cov)

❸ LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                     intercept_scaling=1, l1_ratio=None, max_iter=100,
                     multi_class='auto', n_jobs=None, penalty='l2',
                     random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                     warm_start=False)

[ ] y_predicted_lr = LogRegression.predict(X_test)

[ ] LogRegression.score(X_test,y_test)

❹ 0.9606299212598425

[ ] classification_report(y_test, y_predicted_lr, output_dict = True)

❺ {'0': {'f1-score': 0.9799196787148594,
         'precision': 0.9606299212598425,
         'recall': 1.0,
         'support': 122},
    '1': {'f1-score': 0.0, 'precision': 0.0, 'recall': 0.0, 'support': 5},
    'accuracy': 0.9606299212598425,
    'macro avg': {'f1-score': 0.4899598393574297,
                  'precision': 0.48031496062992124,
                  'recall': 0.5,
                  'support': 127},
    'weighted avg': {'f1-score': 0.9413401638048255,
                     'precision': 0.9228098456196913,
                     'recall': 0.9606299212598425,
                     'support': 127}}
```

SVM

```
[ ] from sklearn.svm import SVC
```

```
[ ] SVC_ = SVC()
```

```
[ ] SVC_.fit(X_cov,y_cov)
```

```
❸ SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
      decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
      max_iter=-1, probability=False, random_state=None, shrinking=True,
      tol=0.001, verbose=False)
```

```
[ ] y_predicted_svc = SVC_.predict(X_test)
y_predicted_svc
```

```
❸ array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
[ ] SVC_.score(X_test,y_test)
```

```
❸ 0.9606299212598425
```

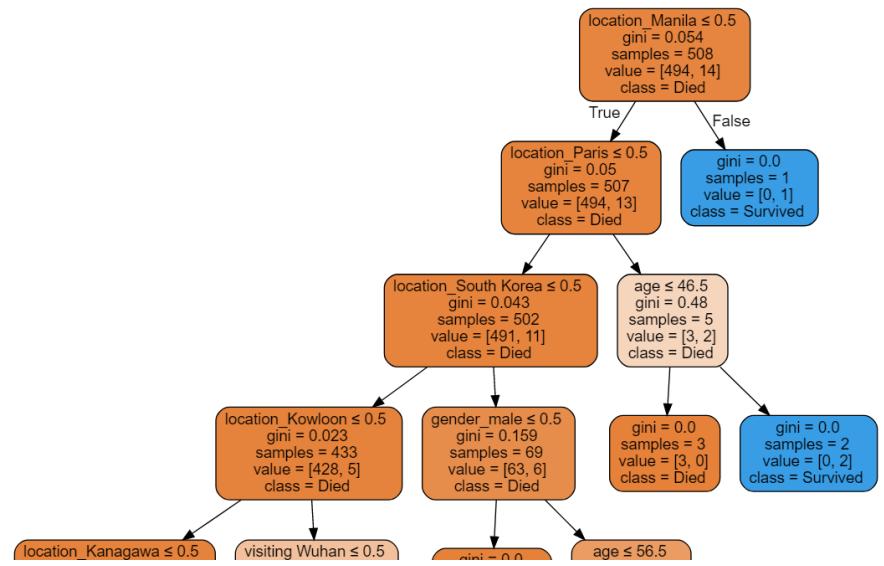
```
[ ] classification_report(y_test,y_predicted_svc, output_dict= True)
```

```
❸ {'0': {'f1-score': 0.9799196787148594,
      'precision': 0.9606299212598425,
      'recall': 1.0,
      'support': 122},
 '1': {'f1-score': 0.0, 'precision': 0.0, 'recall': 0.0, 'support': 5},
 'accuracy': 0.9606299212598425,
 'macro avg': {'f1-score': 0.4899598393574297,
      'precision': 0.48031496062992124,
      'recall': 0.5,
      'support': 127},
 'weighted avg': {'f1-score': 0.9413401638048255,
      'precision': 0.9228098456196913,
      'recall': 0.9606299212598425,
      'support': 127}}
```

Дерево решений

```
[ ] from sklearn.tree import DecisionTreeClassifier  
  
[ ] DTClassifier = DecisionTreeClassifier(random_state=1)  
  
[ ] DTClassifier.fit(X_cov,y_cov)  
  
👤 DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',  
    max_depth=None, max_features=None, max_leaf_nodes=None,  
    min_impurity_decrease=0.0, min_impurity_split=None,  
    min_samples_leaf=1, min_samples_split=2,  
    min_weight_fraction_leaf=0.0, presort='deprecated',  
    random_state=1, splitter='best')  
  
[ ] y_predicted_DT = DTClassifier.predict(X_test)  
  
[ ] DTClassifier.score(X_test,y_test)  
👤 0.937007874015748  
  
[ ] classification_report(y_predicted_DT,y_test, output_dict= True)  
  
👤 {'0': {'f1-score': 0.9674796747967479,  
    'precision': 0.9754098360655737,  
    'recall': 0.9596774193548387,  
    'support': 124},  
    '1': {'f1-score': 0.0, 'precision': 0.0, 'recall': 0.0, 'support': 3},  
    'accuracy': 0.937007874015748,  
    'macro avg': {'f1-score': 0.48373983739837395,  
    'precision': 0.48770491803278687,  
    'recall': 0.4798387096774194,  
    'support': 127},  
    'weighted avg': {'f1-score': 0.9446258242109979,  
    'precision': 0.9523686588356783,  
    'recall': 0.937007874015748,  
    'support': 127}}  
  
[ ] from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, export_graphviz  
    import graphviz
```

```
[ ] #Визуализация дерева решений
dot_data = export_graphviz(DTClassifier, out_file=None,
                           feature_names=data.columns,
                           class_names=('Died','Survived'),
                           filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data)
graph
```



```
[ ] #Важность признаков
```

```
from operator import itemgetter
importance = list(zip(data.columns,DTClassifier.feature_importances_))
importance_sort = sorted(importance, key=itemgetter(1), reverse = True)
importance_sort
```



```
[('age', 0.5623511348413235),
 ('gender_male', 0.16556680750272187),
 ('location_Manila', 0.07612018977332659),
 ('location_Paris', 0.05685496285810745),
 ('visiting_Wuhan', 0.053558249868213),
 ('location_South_Korea', 0.02718938067944169),
 ('location_Kowloon', 0.025130006042144635),
 ('location_Hong_Kong', 0.015302357105203728),
 ('location_Kanagawa', 0.010196776008697743),
 ('location_Japan', 0.004814773664681096),
 ('country_Taiwan', 0.0029153616561387596),
 ('from_Wuhan', 0.0),
 ('recovered', 0.0),
 ('location_Aichi_Prefecture', 0.0),
 ('location_Amiens', 0.0),
 ('location_Andalusia', 0.0),
 ('location_Annecy', 0.0),
 ('location_Baden-Wuerttemberg', 0.0),
 ('location_Barcelona', 0.0),
 ('location_Bavaria', 0.0),
 ('location_Bern', 0.0),
 ('location_Bordeaux', 0.0),
 ('location_Brest', 0.0),
 ('location_California', 0.0),
```

```
[ ] DTClassifier2 = DecisionTreeClassifier(random_state=1)
DTClassifier2.fit(X_cov[['age']],y_cov)
DTClassifier2.score(X_test[['age']],y_test)
```

👤 0.9606299212598425

```
[ ] DTClassifier2 = DecisionTreeClassifier(random_state=1)
DTClassifier2.fit(X_cov[['gender_male']],y_cov)
DTClassifier2.score(X_test[['gender_male']],y_test)
```

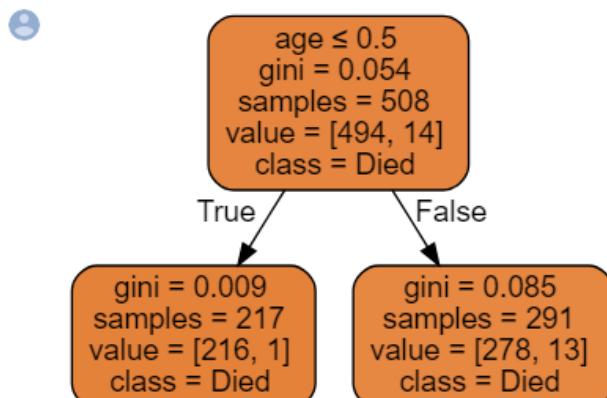
👤 0.9606299212598425

```
[ ] y_predicted_DT2 = DTClassifier2.predict(X_test[['age']])
classification_report(y_predicted_DT2,y_test, output_dict= True)
```

👤

```
{'0': {'f1-score': 0.9799196787148594,
'precision': 1.0,
'recall': 0.9606299212598425,
'support': 127},
'1': {'f1-score': 0.0, 'precision': 0.0, 'recall': 0.0, 'support': 0},
'accuracy': 0.9606299212598425,
'macro avg': {'f1-score': 0.4899598393574297,
'precision': 0.5,
'recall': 0.48031496062992124,
'support': 127},
'weighted avg': {'f1-score': 0.9799196787148594,
'precision': 1.0,
'recall': 0.9606299212598425,
'support': 127}}
```

```
[ ] dot_data1 = export_graphviz(DTClassifier2, out_file=None,
                               feature_names=['age'],
                               class_names=('Died', 'Survived'),
                               filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data1)
graph
```



Лучшими моделями являются логическая регрессия и SVM. Показатели их точности одинаковы.